

Chapter 3

Memory Hierarchy



- **Technology Trend and Memory Hierarchy**
- **Four Questions for Cache Designers**
- **Memory System Performance**
 - **Improve Cache Performance**
- **CPU vulnerabilities**

Memory

- Register
- Cache
- Memory
- Storage
- Mechanical memory
 - Acoustic wave/torque wave delay line memory
 - Magnetic Drum Memory
 - Magnetic core memory
- Electronic memory
 - SRAM
 - DRAM
 - SDRAM
 - Flash
 - ROM
 - PROM
 - EPROM
- Optical memory



Temporal locality (locality in time): if an item is referenced, it will tend to be referenced again soon. If you recently brought a book to your desk to look at, you will probably need to look at it again soon.

Spatial locality (locality in space): if an item is referenced, items whose addresses are close by will tend to be referenced soon.


Memory?

- Load $R2, 0(R1)$
- Store $R2, 0(R1)$




Speed	Processor	Size	Cost (\$/bit)	Current technology
Fastest	Memory	Smallest	Highest	SRAM
	Memory			DRAM
Slowest	Memory	Biggest	Lowest	Magnetic disk


Data processing & temporary storage




Temporary storage




Faster temporary storage



Faster temporary storage



Memory Hierarchy




Memory hierarchy for a personal mobile device

*1000 picoseconds = 1 nanosecond = 10^{-6} millisecond




Memory Hierarchy




Memory hierarchy for a laptop or a desktop

Memory Hierarchy



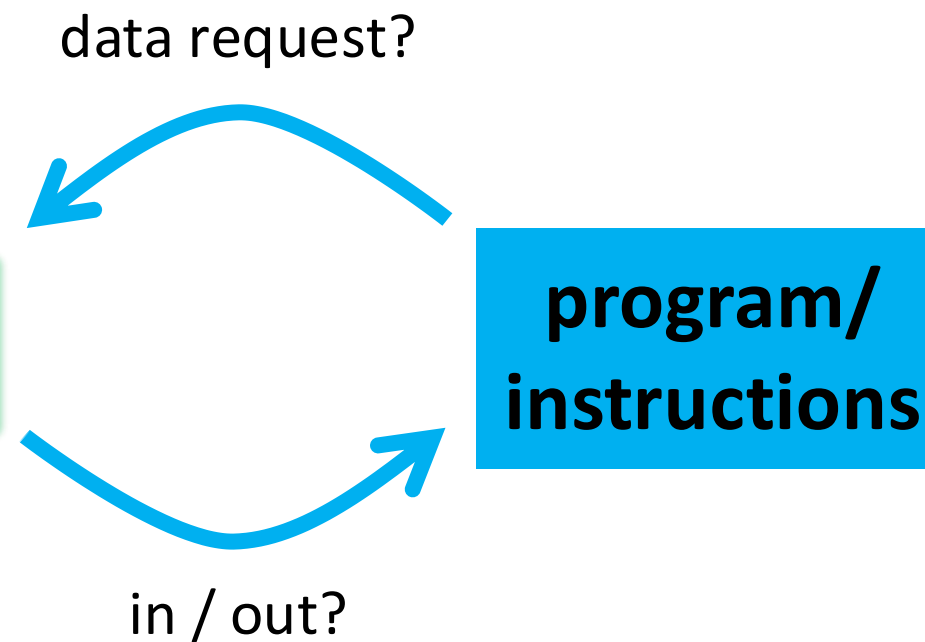
Memory hierarchy for server



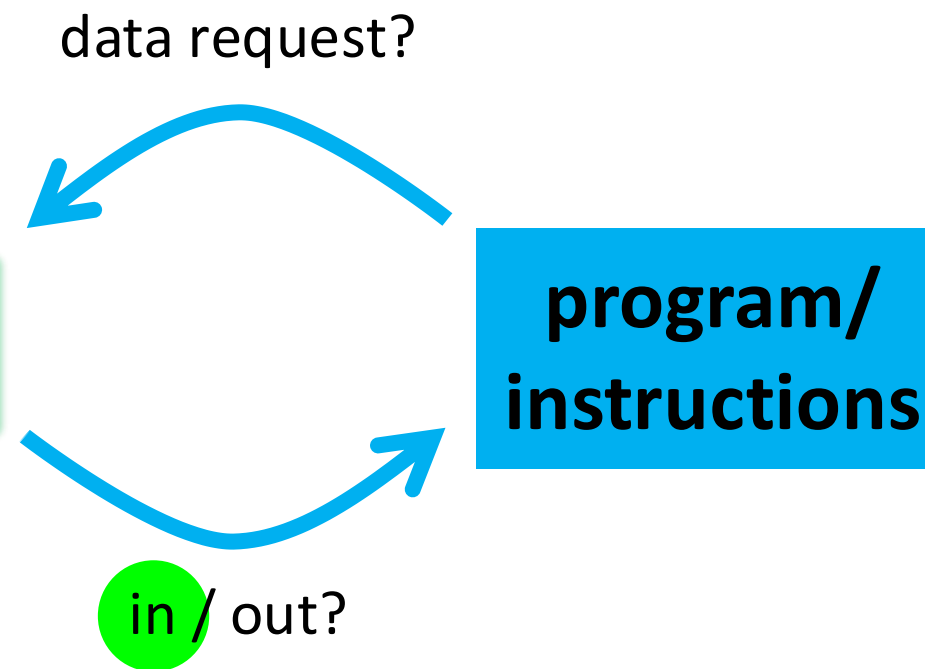
Wait, but what's cache?




Wait, but what's **cache**?



Wait, but what's **cache**?



Wait, but what's **cache**?



So, what's cache?




So, what's cache?

- **Cache: a safe place for hiding or storing things.**

**Webster's New World Dictionary of the
American Language**
Second College Edition (1976)



Cache




- The highest or first level of the memory hierarchy encountered once the *addr* leaves the processor
- Employ buffering to reuse commonly occurring items



Cache Hit/Miss

- When the processor **can/cannot** find a requested data item in the cache



hit rate

miss rate


hit time

miss penalty





Block/Line Run

- A fixed-size collection of data containing the requested word, retrieved from the main memory and placed into the cache




Block/Line Run

- A fixed-size collection of data containing the requested word, retrieved from the main memory and placed into the cache




Block/Line Run

- A fixed-size collection of data containing the requested word, retrieved from the main memory and placed into the cache



Technology Trend



§ 3.2 Technology Trend and Memory Hierarchy

Memory technology	Typical access time	\$ per GiB in 2012
SRAM semiconductor memory	0.5–2.5 ns	\$500–\$1000
DRAM semiconductor memory	50–70 ns	\$10–\$20
Flash semiconductor memory	5,000–50,000 ns	\$0.75–\$1.00
Magnetic disk	5,000,000–20,000,000 ns	\$0.05–\$0.10

SRAM (*static random access memory*)

DRAM (*dynamic random access memory*)

SDRAM (Synchronous DRAM)

Double Data Rate (DDR) SDRAM


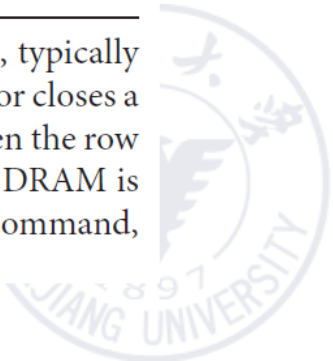



FIGURE 5.4 Internal organization of a DRAM. Modern DRAMs are organized in banks, typically four for DDR3. Each bank consists of a series of rows. Sending a PRE (precharge) command opens or closes a bank. A row address is sent with an Act (activate), which causes the row to transfer to a buffer. When the row is in the buffer, it can be transferred by successive column addresses at whatever the width of the DRAM is (typically 4, 8, or 16 bits in DDR3) or by specifying a block transfer and the starting address. Each command, as well as block transfers, is synchronized with a clock.

HBM (*High Bandwidth Memory*)




Memory/Storage



Flash Memory: electrically erasable programmable read-only memory (EEPROM).




Processor-Memory Performance Gap



Memory Hierarchy of a Modern Computer System

By taking advantage of the principle of locality:

- Present the user with as much memory as is available in the cheapest technology.
- Provide access at the speed offered by the fastest technology.



Cache Locality

- **Temporal locality**

need the requested word again soon

- **Spatial locality**

likely need other data within the same block soon



Cache Miss

- Time required for cache miss depends on:
 - **Latency**: the time to retrieve the first word of the block
 - **Bandwidth**: the time to retrieve the rest of this block
- Causes of Cache misses
 - **Compulsory**: first reference to a block
 - **Capacity**: blocks discarded and later retrieved
 - **Conflict**: program makes repeated references to multiple addresses from different blocks that map to the same location in the cache



Enhance performance

By taking advantage of the principle of locality:

- **most programs do not access all code or data uniformly**
- **Temporal Locality** (Locality in Time):
 - If an item is referenced, the **same item** will tend to be referenced again **soon**
 - Keep most recently accessed data items closer to the processor
- **Spatial Locality** (Locality in Space):
 - If an item is referenced, **nearby items** will tend to be referenced **soon**
 - Move recently accessed groups of contiguous words(block) closer to processor.



Three classes of computers with different concerns in memory hierarchy

- **Desktop computers:**

- Are primarily running one application for single user
- Are concerned more with average latency from the memory hierarchy.

- **Server computers:**

- May typically have hundreds of users running potentially dozens of applications simultaneously.
- Are concerned about memory bandwidth.



Three classes of computers with different concerns in memory hierarchy

- **Embedded computers:**
 - Real-time applications.
 - Worst-case performance vs Best case performance
 - Are concerned more about power and battery life.
 - Hardware vs software
 - Running single app & use simple OS
 - The protection role of the memory hierarchy is often diminished.
 - Main memory is very small
 - often no disk storage




36 terms of Cache

Cache	full associative	write allocate
Virtual memory	dirty bit	unified cache
Memory stall cycles	block	block offset
misses per instruction	direct mapped	write back
Valid bit	data cache	locality
Block address	hit time	address trace
Write through	cache miss	set
Instruction cache	page fault	miss rate
random replacement	index field	cache hit
Average memory access time	page	tag field
n-way set associative	no-write allocate	miss penalty
Least-recently used	write buffer	write stall





What is a cache?

- Small, fast storage used to improve average access time to slow memory.
- In computer architecture, almost everything is a cache!
 - Registers “a cache” on variables – software managed
 - First-level cache a cache on second-level cache
 - Second-level cache a cache on memory
 - Memory a cache on disk (virtual memory)
 - TLB a cache on page table
 - Branch-prediction a cache on prediction information?



§ 3.2 Technology Trend and Memory Hierarchy






A typical multi-level cache organization

Split vs. unified caches

- Unified cache
 - All memory requests go through a single cache.
 - This requires less hardware, but also has lower performance
- Split I & D cache
 - A separate cache is used for instructions and data.
 - This uses additional hardware, though there are some simplifications (the I cache is read-only).



Cache type	Associativity (E)	Block size (B)	Sets (S)	Cache size (C)
on-chip L1 i-cache	4	32 B	128	16 KB
on-chip L1 d-cache	4	32 B	128	16 KB
off-chip L2 unified cache	4	32 B	1024–16384	128 KB–2 MB

Intel Pentium cache organization

The cache just before and just after a reference to a word X_n that is not initially in the cache.

X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_3

a. Before the reference to X_n


X_4
X_1
X_{n-2}
X_{n-1}
X_2
X_n
X_3


b. After the reference to X_n

This reference causes a miss that forces the cache to fetch X_n from memory and insert into the cache.

How do we know if a data item is in the cache?
Moreover, if it is, how do we find it?







Lines (i.e., blocks) of data are copied from memory to the cache.

Memory Access **without** Cache

- Load word instruction: `lw t0, 0(t1)`
- `t1` contains 1022_{10} , `Memory[1022]=99`
 - Processor issues address 1022_{10} to Memory
 - Memory reads word at address $1022_{10}(99)$
 - Memory sends `99` to Processor
 - Processor loads `99` into register `t0`

Memory Access **with** Cache

- Load word instruction: `lw t0, 0(t1)`
- `t1` contains 1022_{10} , `Memory[1022]=99`
- With cache (similar to a hash)
 - Processor issues address 1022_{10} to Cache
 - Cache checks to see if has copy of data at address 1022_{10}
 - If finds a match (**cache hit**): cache reads `99`, sends to process
 - Not match (**cache miss**): cache sends address to 1022_{10} Memory
 - Memory reads `99` at address 1022_{10}
 - Memory sends `99` to Cache
 - Cache replaces word with **new 99**
 - Cache sends `99` into process
 - Processor loads `99` into register `t0`

Memory Access **with** Cache

- When reading memory, three things can happen:
- **Cache hit**
 - Cache block is valid and contains proper address, so read desired word
- **Cache miss**
 - Nothing in cache in appropriate block, so fetch from memory
- **Cache miss, block replacement**
 - Wrong data is in cache at appropriate block, so discard it and fetch desired data from memory (cache always copy)

Four Questions for Cache Designers

Caching is a general concept used in processors, operating systems, file systems, and applications.

There are **Four Questions** for Cache/Memory Hierarchy Designers


- **Q1:** Where can a block be placed in **the upper level/main memory**?
(Block placement)
 - Fully Associative, Set Associative, Direct Mapped
- **Q2:** How is a block found if it is in **the upper level/main memory**?
(Block identification)
 - Tag/Block
- **Q3:** Which block should be replaced on a **Cache/main memory** miss?
(Block replacement)
 - Random, LRU, FIFO
- **Q4:** What happens on a write?
(Write strategy)
 - Write Back or Write Through (with Write Buffer)




Q1: Block Placement

Direct mapped

(Block address) modulo (Number of blocks in the cache)




A cache structure in which each memory location is mapped to exactly one location in the cache.




Q1: Block Placement

Direct mapped



(Block address) modulo (Number of blocks in the cache)



A cache structure in which each memory location is mapped to exactly one location in the cache.





Fully-associative



A cache structure in which a block can be placed in any location in the cache.




2-way Set-associative




A cache that has a fixed number of locations (at least two) where each block can be placed.

4-way Set-associative



Why Index With the Middle Bits?

If the high-order bits are used as an index, then some contiguous memory blocks will map to the same cache set.




Q1: Block Placement

- Direct mapped
 - Block
 - *Note that direct mapped is the same as 1-way set associative, and fully associative is m-way set-associative (for a cache with m blocks).*
- Fully associative
- Set associative
 - Block can go in one of a set of places in the cache.
 - A set is a group of blocks in the cache.
 - Block address **MOD** Number of sets in the cache
 - If sets have n blocks, the cache is said to be n-way set associative.



n sets => n -way set associative
Direct-mapped cache => one block per set
Fully associative => one set

8-32 Block Placement



N-way Set-associative

- The higher the degree of association, the higher the utilization of cache space, the lower the probability of block collision and the lower the failure rate.

	n	G
Full-associative	M	1
Direct mapped	1	M
Set-associative	$1 < n < M$	$1 < G < M$

- Most Cache: $n \leq 4$
- Question: Is the greater the number n , the better?



Four Questions for Cache Designers

Caching is a general concept used in processors, operating systems, file systems, and applications.

There are **Four Questions** for Cache/Memory Hierarchy Designers

- **Q1:** Where can a block be placed in **the upper level/main memory**?
(Block placement)
 - Fully Associative, Set Associative, Direct Mapped
- **Q2:** How is a block found if it is in **the upper level/main memory**?
(Block identification)
 - Tag/Block
- **Q3:** Which block should be replaced on a **Cache/main memory** miss?
(Block replacement)
 - Random, LRU, FIFO
- **Q4:** What happens on a write?
(Write strategy)
 - Write Back or Write Through (with Write Buffer)





Q2: Block Identification

- Every block has an **address tag** that stores the main memory address of the data stored in the block.
- When checking the cache, the processor will **compare** the requested **memory address to the cache tag** -- if the two are equal, then there is a cache hit and the data is present in the cache
- Often, each cache block also has a **valid bit** that tells if the contents of the cache block are valid




The Format of the Physical Address

- The **Index** field selects
 - The **set**, in case of a **set-associative cache**
 - The **block**, in case of a **direct-mapped cache**
 - Has as many bits as $\log_2(\text{\#sets})$ for **set-associative caches**, or $\log_2(\text{\#blocks})$ for **direct-mapped caches**
- The **Byte Offset** field selects
 - The byte within the block
 - Has as many bits as $\log_2(\text{size of block})$
- The **Tag** is used to find the matching block within a set or in the cache
 - Has as many bits as $\text{Address_size} - \text{Index_size} - \text{Byte_Offset_Size}$



Line Matching and Word Selection




direct-mapped cache

Direct-mapped Cache Example (1-word Blocks)


LOAD R1, 0x04

MEMORY

Address	Data
0x00	0x00000000
0x04	0x12345678
0x08	0x87654321
0x0C	0x11111111
0x10	0x22222222
0x14	0x33333333
0x18	0x44444444
0x1C	0x55555555
0x20	0x10101010




Line Matching and Word Selection




set associative cache

Fully-Associative Cache example (1-word Blocks)

- Assume cache has 4 blocks




Line Matching and Word Selection



fully associative cache


2-Way Set-Associative Cache

- Assume cache has 4 blocks and each block is 1 word
- 2 blocks per set, hence 2 sets per cache



What is a cache?

- Small, fast storage used to improve average access time to slow memory.
- In computer architecture, almost everything is a cache!
 - Registers “a cache” on variables – software managed
 - First-level cache a cache on second-level cache
 - Second-level cache a cache on memory
 - Memory a cache on disk (virtual memory)
 - TLB a cache on page table
 - Branch-prediction a cache on prediction information?




Using a Finite-State Machine to Control a Simple Cache

finite-state machine


A sequential logic function consisting of a set of inputs and outputs, next-state function that maps the current state and the inputs to a new state, and an output function that maps the current state and possibly the inputs to a set of asserted outputs.

next-state function

A combinational function that, given the inputs and the current state, determines the next state of a finite-state machine.



Using a Finite-State Machine to Control a Simple Cache



Access data in a direct mapped cache

- 16KB of data, direct-mapped, 4 words blocks
- Read 4 addresses:
 - 0x00000014
 - 0x0000001C
 - 0x00000034
 - 0x00008014
- Memory values here:

Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- 16KB of data, direct-mapped, 4 words blocks
- Read 4 addresses:

• 0x00000014	00000000000000000000 0000000001 0100
• 0x0000001C	00000000000000000000 0000000001 1100
• 0x00000034	00000000000000000000 0000000011 0100
• 0x00008014	00000000000000000010 0000000001 0100

- Memory values here:
- 4 addresses divided into Tag, Index, Byte Offset

Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l

Access data in a direct mapped cache

- Valid bit: determines whether anything is stored in that row (when computer initially powered up, all entries invalid)


		Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index	0	0	g				
	1	0					
	2	0					
	3	0					
	4	0					
	5	0					
	6	0					
	7	0					
...		...					
1022	1022	0					
	1023	0					

Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	
00000030	e
00000034	f
00000038	g
0000003C	h
...	
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00000014, read block 1 (0000000001)




Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00000014, no valid data



Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00000014, load that data into cache, setting tag, valid

■ 000000000000000000000000 0000000001 0100
 Tag Index Offset

Index	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	1	g	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...						
1022	0					
1023	0					

Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l

Access data in a direct mapped cache

- Read 0x00000014, read from cache at offset, return word b

00000000000000000000 0000000001 0100


	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index						
0	0					
1	1	g	d	c	b	a
2	0					
3	0					
4	0					
5	0					
6	0					
7	0					
...						
1022	0					
1023	0					

Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x0000001C, index is valid




Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x0000001C, index is valid, tag matches




Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x0000001C, index is valid, tag matches, return d



Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00000034, read block 3, no valid data

▪ 00000000000000000000 0000000011 0100

Tag Index Offset

Index	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
0	0					
1	1	g				
2	0		d	c	b	a
3	0					
4	0					
5	0					
6	0					
7	0					
...						
1022	0					
1023	0					


Memory

Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00000034, load that cache block, return word f



Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00008014, read cache block 1, data is valid, cache block 1 tag does not match (0!=2)



Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00008014, miss, replace block 1 with new data & tag

▪ 00000000000000000010 0000000001 0100

	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index						
0	0					
1	1	g				
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...			...			
1022	0					
1023	0					

Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Access data in a direct mapped cache

- Read 0x00008014, return word j

00000000000000000010 0000000001 0100

	Valid	Tag	0xc-f	0x8-b	0x4-7	0x0-3
Index						
0	0					
1	1	g		k	j	i
2	0					
3	1	0	h	g	f	e
4	0					
5	0					
6	0					
7	0					
...						
1022	0					
1023	0					

Address (hex)	Memory Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l



Homework

Cache 16KB of data, direct-mapped, 4 words blocks Memory

Index						
x	0	0				
	1	1	2	l	k	i
	2	0				
	3	1	0	h	q	f
	4	0				
	5	0				
	6	0				
	7	0				

- (1) Read address 0x00000030, what happens?
- (2) Read address 0x0000001c, what happens?

Please present cache hit or miss, the index, tag match or mismatch, offset, and the return value.

Address (hex)	Value of Word
00000010	a
00000014	b
00000018	c
0000001C	d
...	...
00000030	e
00000034	f
00000038	g
0000003C	h
...	...
00008010	i
00008014	j
00008018	k
0000801C	l

Four Questions for Cache Designers

Caching is a general concept used in processors, operating systems, file systems, and applications.


There are **Four Questions** for Cache/Memory Hierarchy Designers

- **Q1:** Where can a block be placed in **the upper level/main memory**?
(Block placement)
 - Fully Associative, Set Associative, Direct Mapped
- **Q2:** How is a block found if it is in **the upper level/main memory**?
(Block identification)
 - Tag/Block
- **Q3:** Which block should be replaced on a **Cache/main memory** miss?
(Block replacement)
 - Random, LRU, FIFO
- **Q4:** What happens on a write?
(Write strategy)
 - Write Back or Write Through (with Write Buffer)




§ 3.3 Four Questions for Cache Designers


tag: A field in a table used for a memory hierarchy that contains the address information required to identify whether the associated block in the hierarchy corresponds to a requested word.




Direct-mapped cache




Set associative cache




Fully set associative cache



Set selection in a direct-mapped cache




Set selection in a set associative cache



Set selection in a fully associative cache

The size of the block was one word (4 bytes).



- 64-bit addresses
- A direct-mapped cache
- The cache size is 2^n blocks, so n bits are used for the index
- The block size is
 - ✓ 2^m words (2^{m+2} bytes = 2^{m+5} bits)
 - ✓ m bits are used for the word within the block, and two bits are used for the byte part of the address

$$\text{tag} = 64 - (n + m + 2)$$

$$\begin{aligned} \text{cache} &= 2^n \times (\text{block size} + \text{tag size} + \text{Valid size}). \\ &= 2^n \times (2^m \times 32 + (64 - n - m - 2) + 1) \\ &= 2^n \times (2^m \times 32 + 63 - n - m). \end{aligned}$$

actual size/complete cache size



Bits in a Cache

Example 1: How many total bits are required for a direct-mapped cache with 16 KiB of data and four-word blocks, assuming a 64-bit address?

Answers:

16 KiB = 4096 words = 2^{12} words

The block size is 4 words = 2^2 words, so $m=2$

$2^{12} / 2^2 = 1024$ blocks, so $n=10$

Each block has $4 \times 32 = 128$ bits of data plus a tag, plus a valid bit

tag = $64 - (n + m + 2) = 64 - (10 + 2 + 2) = 50$ bits

Cache size = $2^n \times (\text{block size} + \text{tag size} + \text{Valid size})$

$$= 2^{10} \times (4 \times 32 + 50 + 1) = 179 \text{ Kib} = 22.375 \text{ KiB}$$



Mapping an Address to a Multiword Cache Block

Example 2: Consider a cache with 64 blocks and a block size of 16 bytes. To what block number does byte address 1200 map?

Answers:

The cache block number

= (Block address) modulo (Number of blocks in the cache)

Block address = Byte address / Bytes per block = $1200 / 16 = 75$

$75 \text{ modulo } 64 = 11$

In fact, this block maps all addresses between 1200 and 1215.



Four Questions for Cache Designers

Caching is a general concept used in processors, operating systems, file systems, and applications.

There are **Four Questions** for Cache/Memory Hierarchy Designers

- **Q1:** Where can a block be placed in **the upper level/main memory**?
(Block placement)
 - Fully Associative, Set Associative, Direct Mapped
- **Q2:** How is a block found if it is in **the upper level/main memory**?
(Block identification)
 - Tag/Block
- **Q3:** Which block should be replaced on a **Cache/main memory** miss?
(Block replacement)
 - Random, LRU, FIFO
- **Q4:** What happens on a write?
(Write strategy)
 - Write Back or Write Through (with Write Buffer)



Q3: Block Replacement

- In a direct-mapped cache, there is only one block that can be replaced



§ 3.3 Four Questions for Cache Designers

Binary address of reference	Assigned cache block (where found or placed)
10110 _{two}	(10 110 _{two} mod 8) = 110 _{two}
11010 _{two}	(11 010 _{two} mod 8) = 010 _{two}
10110 _{two}	(10 110 _{two} mod 8) = 110 _{two}
11010 _{two}	(11 010 _{two} mod 8) = 010 _{two}
10000 _{two}	(10 000 _{two} mod 8) = 000 _{two}
00011 _{two}	(00 011 _{two} mod 8) = 011 _{two}
10000 _{two}	(10 000 _{two} mod 8) = 000 _{two}
10010 _{two}	(10 010 _{two} mod 8) = 010 _{two}
10000 _{two}	(10 000 _{two} mod 8) = 000 _{two}

Index	V	Tag	Data
000	N		
001	N		
010	N		
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

Index	V	Tag	Data
000	N		
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	N		
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	11 _{two}	Memory (11010 _{two})
011	Y	00 _{two}	Memory (00011 _{two})
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

Index	V	Tag	Data
000	Y	10 _{two}	Memory (10000 _{two})
001	N		
010	Y	10 _{two}	Memory (10010 _{two})
011	Y	00 _{two}	Memory (00011 _{two})
100	N		
101	N		
110	Y	10 _{two}	Memory (10110 _{two})
111	N		

10110

miss

11010

miss

10110

hit

11010

hit

10000

miss

00011

miss

10000

hit

10010

miss

10000

hit



cache hit

If the cache reports a hit, the computer continues using the data as if nothing happened.

cache miss

For a cache miss, we can stall the entire processor, essentially freezing the contents of the temporary and programmer-visible registers, while we wait for memory.




We can now define the steps to be taken on an instruction cache miss:

1. Send the original PC value to the memory. (PC-4)
2. Instruct main memory to perform a read and wait for the memory to complete its access.
3. Write the cache entry, putting the data from memory in the data portion of the entry, writing the upper bits of the address (from the ALU) into the tag field, and turning the valid bit on.
4. Restart the instruction execution at the first step, which will refetch the instruction, this time finding it in the cache.



Q3: Block Replacement

- In a direct-mapped cache, there is only one block that can be replaced
- In set-associative and fully-associative caches, there are N blocks (where N is the degree of associativity)



Strategy of Block Replacement

- Several different replacement policies can be used
 - **Random replacement** - *randomly pick any block*
 - Easy to implement in hardware, just requires a random number generator
 - Spreads allocation uniformly across cache
 - May evict a block that is about to be accessed
 - **Least-Recently Used (LRU)** - *pick the block in the set which was least recently accessed*
 - Assumed more recently accessed blocks more likely to be referenced again
 - This requires extra bits in the cache to keep track of accesses.
 - **First In, First Out (FIFO)** - *Choose a block from the set which was first came into the cache*



Strategy of Block Replacement

Suppose:


- Cache block size is 3, and access sequence is shown as follows.

2, 3, 2, 1, 5, 2, 4, 5, 3, 4

- FIFO, LRU and OPT are used to simulate the use and replacement of cache block.



FIFO




LRU

		↓	↓	↓	↓	↓	↓	↓	↓	↓	
Time		1	2	3	4	5	6	7	8	9	10
Sequence		2	3	2	1	5	2	4	5	3	4
3 5 4		2	2 3	2 3 hit	2 3* 1	2* 5 1	2 5 1* hit	2 5* 4	2* 5 4 hit	3 5 4*	3 5* 4 hit



OPT



The hit rate is related to the replacement algorithm.



Time	1	2	3	4	5	6	7	8	thrashing
Sequence	1	2	3	4	1	2	3	4	
FIFO	1	1	1*	4	4	4*	3	3	No hit
n=3		2	2	2*	1	1	1*	4	
			3	3	3*	2	2	2*	
LRU	1	1	1*	4	4	4*	1	3	Hit 4 times
n=4		2	2	2*	2	2	2*	2	
OPT			3	3	3*	3	3	3*	Hit 3 times
	1	1	1	4	4*	4	4	4	
		2	2	2	2	2*	3*	3	
			3*	4*	4	4	4	4*	

➤ Hit rate is related to access sequence.



Stack replacement algorithm

- $B_t(n)$ represents the set of access sequences contained in a cache block of size n at time t .
- $B_t(n)$ is the subset of $B_t(n + 1)$.



LRU replacement algorithm is stack replacement algorithm


Time	1	2	3	4	5	6	7	8	9	10	11	12
Sequence	1	2	3	4	1	2	5	1	2	3	4	5
LRU n=3 Bt(n)	1	1 2	1* 2 3	4 2* 3	4 1 3*	4* 1 2	5 1* 2	5 1 2*	5* 1 2	3 1* 2	3 4 2*	3* 4 5
LRU n+1=4 Bt(n+1)	1	1 2	1 2 3	1* 2 3 4	1 2* 3 4	1 2 3* 4	1 2 5 4*	1 2 5 4*	1 2 5 4*	1 2 5* 3	1* 2 5* 4	5 2* 4 3




Time	1	2	3	4	5	6	7	8	9	10	11	12
Sequence	1	2	3	4	1	2	5	1	2	3	4	5
FIFO n=3	1	1 2	1* 2 3	4 2* 3	4 1 3*	4* 1 2	5 1* 2	5 1* 2	5 1* 2	5 3 2*	5* 3 4	5* 3 4
FIFO n+1=4	1	1 2	1 2 3	1* 2 3 4	1* 2 3 4	1* 2 3 4	5 2* 3 4	5 1 3* 4	5 1 2 4*	5* 1 2 3	4 1* 2 3	4 5 2* 3



Using LRU





For LRU algorithm, the hit ratio always increases with the increase of cache block.

Time	1	2	3	4	5	6	7	8	9	10	11	12
Sequence	1	2	3	4	1	2	5	1	2	3	4	5
LRU n=3 Hit 2 times	1	1 2	1* 2 3	4 2* 3	4 1 3*	4* 1 2	5 1* 2	5 1 2*	5* 1 2	3 1* 2	3 4 2*	3* 4 5
LRU n=4 Hit 4 times	1	1 2	1 2 3	1* 2 3 4	1 2* 3 4	1 2 3* 4	1 2 5 4*	1 2 5 4*	1 2 5 4*	1 2 5* 3	1* 2 4 3	5 2* 4 3
					Hit	Hit		Hit	Hit			



Belady

Time	1	2	3	4	5	6	7	8	9	10	11	12	
Sequence	1	2	3	4	1	2	5	1	2	3	4	5	
FIFO n=3 Hit 3 times	1	1 2	1* 2 3	4 2* 3	4 1 3*	4* 1 2	5 1* 2	5 1* 2	5 1* 2	5 3 2*	5* 3 4	5* 3 4	5* 3 4
								Hit	Hit			Hit	
FIFO n+1=4 Hit 2 times	1	1 2	1 2 3	1* 2 3 4	1* 2 3 4	1* 2 3 4	5 2* 3 4	5 1 3* 4	5 1 2 4*	5* 1 2 3	5* 1* 2 3	4 1* 2 3	4 5 2* 3
					Hit	Hit							



LRU

- How can I implement the LRU replacement algorithm with only ordinary gates and triggers?
- Comparison Pair Method
- Basic idea: Let each cache block be combined in pairs, use a **comparison pair flip-flop** to record the order in which the two cache blocks have been accessed in the comparison pair, and then use a gate circuit to combine the state of each comparison pair flip-flop, you can find the block to be replaced according to the LRU algorithm.



Example

- There are three cache blocks (A, B, and C), which can be combined into 6 pairs (AB, BA, AC, CA, BC, and CB). Among them, AB and BA, AC and CA, BC and CB are repeated, so only take AB, AC, BC.
- The access sequence of each pair is represented by “comparison pair flip-flops” T_{AB} , T_{AC} , and T_{BC} respectively. T_{AB} is "1", which means that A has been accessed more recently than B; T_{AB} is "0", which means that B has been accessed more recently than A. T_{AC} and T_{BC} are similarly defined.




- If the most recently accessed block is A and C is the block that has not been accessed for the longest time, the three flip-flops' states must be respectively: $T_{AB}=1$, $T_{AC}=1$, and $T_{BC}=1$.
- If the most recently accessed block is B and C is the block that has not been accessed for the longest time, the three flip-flops' states must be respectively: $T_{AB}=0$, $T_{AC}=1$, and $T_{BC}=1$.
- Therefore, the block C that has not been accessed for the longest will be replaced. In that, the Boolean algebra expression must be:

$$C_{LRU} = T_{AB} \bullet T_{AC} \bullet T_{BC} + \overline{T_{AB}} \bullet T_{AC} \bullet T_{BC} = T_{AC} \bullet T_{BC}$$

$$B_{LRU} = T_{AB} \bullet \overline{T_{BC}}$$

$$A_{LRU} = \overline{T_{AB}} \bullet \overline{T_{AC}}$$





- Change the state of the flip-flop after each access.
 - After accessing block A: $T_{AB}=1$, $T_{AC}=1$
 - After accessing block B: $T_{AB}=0$, $T_{BC}=1$
 - After accessing block C: $T_{AC}=0$, $T_{BC}=0$



Hardware usage analysis (if p is the number of cache blocks)

- Since each block may be replaced, its signal needs to be generated with an AND gate, so the number of AND gates will be equal to p .
- Each AND gate receives inputs from its related flip-flops, for example, A_{LRU} AND gates must have inputs from T_{AB} and T_{AC} , B_{LRU} must have inputs from T_{AB} and T_{BC} , and the number of comparison pair flip-flops is the block number minus 1, so the input number of the AND gate is $p - 1$.
- If p is the block number, for pairwise combination, the number of comparison pair flip-flops should be C_p^2 , which is $p \cdot (p-1)/2$.



The Relationship between the Block Number and Required Hardware for Comparison Pair

Block Number	3	4	6	8	16	64	256
Number of Flip-flop	3	6	15	28	120	2016	32640
Number of And-Gate	3	4	6	8	16	64	256
Input Number of And-Gate	2	3	5	7	15	63	255



Homework

For a direct-mapped cache design with a 64-bit address, the following bits of the address are used to access the cache.

Tag	Index	Offset
63-10	9-5	4-0

- (1) What is the cache block size (in words)?
- (2) How many blocks does the cache have?
- (3) What is the ratio between total bits required for such a cache implementation over the data storage bits?

Beginning from the power on, the following byte-addressed cache references are recorded.

Address												
Hex	00	04	10	84	E8	A0	400	1E	8C	C1C	B4	884
Dec	0	4	16	132	232	160	1024	30	140	3100	180	2180

- (1) For each reference, list i) its tag, index, and offset, ii) whether it is a hit or a miss, and iii) which bytes were replaced (if any).
- (2) What is the hit ratio?

